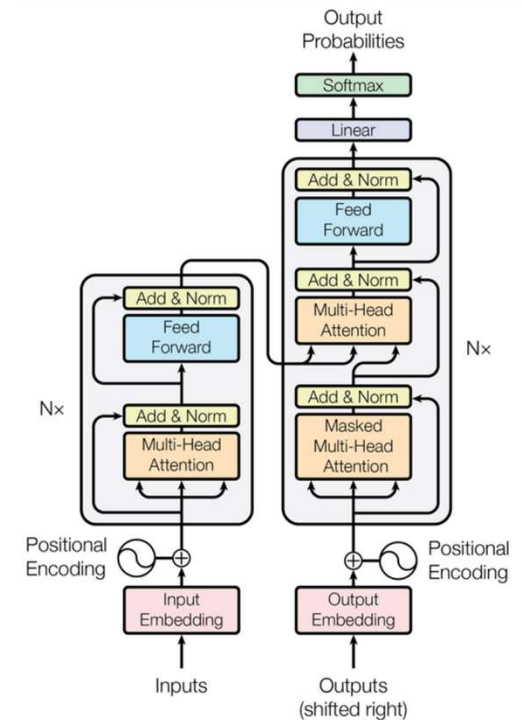


Introduction to

- (1) AI-Assisted Coding
- (2) Autoregressive Large Language Model (LLM)
- (3) Prompt Engineering 201

Minha Hwang





AI-ASSITED CODING DEMO

- NOTEBOOK: *AI_ASSITED_CODING.IPYNB*

LANGUAGE MODEL (LM) - AUTOREGRESSIVE

LM (language model): predict next word given input

- Input: text
- Output: next word prediction



Image source: <https://web.Stanford.edu/class/cs224n>

AUTOREGRESSIVE LANGUAGE MODELS: PROBABILITY KERNEL

- A language model is a **probability kernel μ** given a prefix of words: **$\mu: X \rightarrow Pr(Y)$**
 - Stochastic in nature: **A same prefix X** can give a **random output** sampled from a probability distribution μ_X (i.e., generative) → A key reason for factual inaccuracy, inconsistency or hallucination (making stuff up)
- A language model calculates **$Pr(s)$** given a sequence of words: **$s = (w_1, w_2, \dots, w_{T-1}, w_T)$**
- An autoregressive language model calculates this **conditional on a previous sequence of words**:

$$\begin{aligned} Pr(s) &= Pr(w_1, w_2, \dots, w_{T-1}, w_T) \\ &= \prod_{t=1}^T Pr(w_t | w_1, w_2, \dots, w_{t-1}) \end{aligned}$$

- **Next-word prediction**: Given a prefix $(w_1, w_2, \dots, w_{t-1})$, calculate the probability of the next word w_t (Conceptually same to time series with path dependence)

Source: Prof. Kyunghyun Cho

AUTOREGRESSIVE LANGUAGE MODELS: SIMPLE EXAMPLE

- 4-word sentence example: “I am a student”

$$Pr(s) = Pr(w_1, w_2, w_3, w_4) = Pr(w_1) \times Pr(w_2|w_1) \times Pr(w_3|w_1, w_2) \times Pr(w_4|w_1, w_2, w_3)$$

- All you need is “counting” (if there are large amounts of data)

$$Pr(w_2|w_1) = \frac{count(w_1, w_2)}{count(w_1)} \longrightarrow \text{2-grams (Bigrams)}$$

$$Pr(w_3|w_1, w_2) = \frac{count(w_1, w_2, w_3)}{count(w_1, w_2)} \longrightarrow \text{3-grams (Trigrams)}$$

$$Pr(w_4|w_1, w_2, w_3) = \frac{count(w_1, w_2, w_3, w_4)}{count(w_1, w_2, w_3)} \longrightarrow \text{4-grams}$$

- Problems:
 - This requires **a lot of space (RAM)**
 - Count-based language models **cannot generalize**: A certain sentence **does not appear** in the corpus

Source: Prof. Kyunghyun Cho

TWO TYPES OF LARGE LANGUAGE MODEL (LLM)

Base LLM: Pre-Training

predict next word, based on text training data

- Self-supervised

Once upon a time, there was a unicorn
that lived in a magical forest with
all her unicorn friends

What is the capital of France?

What is France's largest city?

What is France's population?

What is the currency of France?

Instruction Tuned LLM: Post Training

Tries to follow instructions

Fine-tune on instructions and good attempts at following those instructions (SFT)

- Human Labeled Data: Instruction – Response Pair
- RLHF (Reinforcement Learning with Human Feedback) or DPO (Direct Preference Optimization)

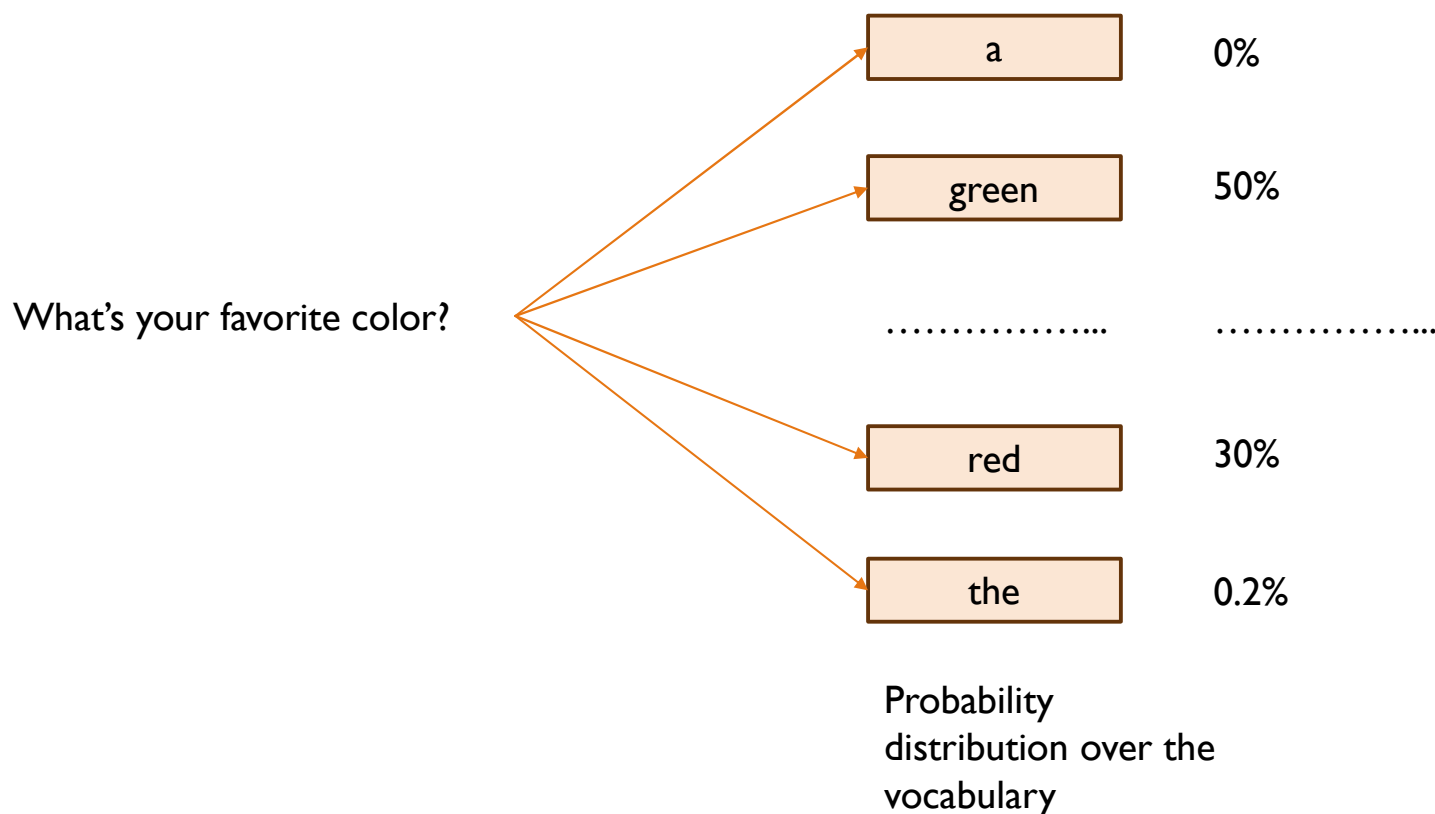
What is the capital of France?

The capital of France is Paris.

Helpful, Honest, Harmless (Style)

Source: DeepLearning.AI – ChatGPT Prompt Engineering

SAMPLING: KEY TO UNDERSTAND PARAMETERS FOR LLM

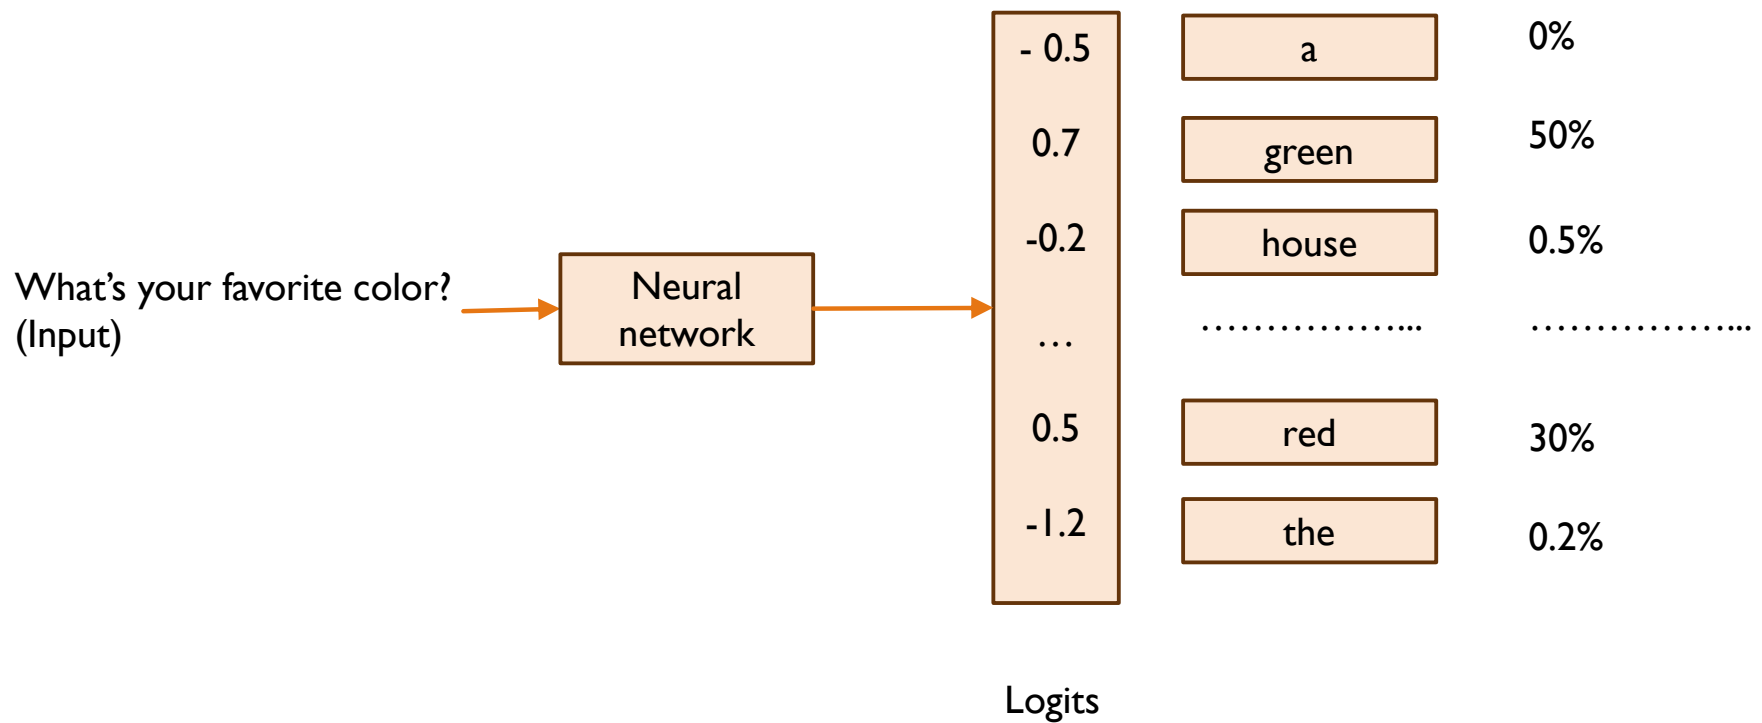


SoftMax:

Multiple classes

- **Greedy sampling:**
Always pick the outcome with highest probability (boring outcome)
- Sample the next token **according to the probability distribution** over all possible values

UNDER THE HOOD



Source: AI Engineering

TEMPERATURE: CONTROLS RANDOMNESS/CREATIVITY (1/2)

- A constant used to adjust the logits vector before the softmax transformation
- Logits are divided by temperature
- Higher temperature reduces the probabilities of common tokens: increase the probabilities of rarer tokens
 - (+) More creative model responses
 - (-) More hallucination, inconsistency

1. Adjusting Logits with Temperature

Given a logit vector z , temperature T scales the logits as follows:

$$z' = \frac{z}{T}$$

where:

- z is the original logit vector output from the model.
- T is the **temperature** hyperparameter.
- z' is the adjusted logit vector.
- If $T = 1$, logits remain unchanged.
- If $T > 1$, logits are **flattened** (making the probability distribution more uniform).
- If $T < 1$, logits are **sharpened** (making high probabilities even higher and low probabilities even lower).

2. Applying Softmax to Adjusted Logits

Once we have the temperature-scaled logits z' , we apply the **softmax function**:

$$P_i = \frac{e^{z'_i}}{\sum_j e^{z'_j}}$$

where:

- P_i is the probability of token i .
- $z'_i = \frac{z_i}{T}$ is the temperature-scaled logit.
- The denominator ensures that the probabilities sum to 1.

TEMPERATURE (2/2)



```
{  
  "prompt": "Give me a list of 5 unusual ice cream flavors.",  
  "max_tokens": 30,  
  "temperature": 0.7  
}
```

Example:

- With temperature = 0.7, you may get some unique, creative flavors like “Lavender Basil” or “Sriracha Peanut.”
- If you set temperature = 0.0, the model’s answer might be more “plain” or it might return more common “unique” flavors, like “Mint Chocolate Chip.”

What it is:

- A value between 0 and 2 (commonly 0 to 1) that controls the “creativity” or randomness in the model’s outputs.

Why it matters:

- **Lower temperature** (e.g., 0–0.3): The output is **more deterministic and focused**, which is helpful for tasks requiring **factual or logical consistency** (e.g., summarization, factual Q&A).
- **Higher temperature** (e.g., 0.7–1): The output becomes **more creative**, with **increased risk of straying off-topic** or introducing **less relevant details**.

LOGPROBS: LOG PROBABILITIES

What it is

- Log Probs: Probability in log scale (helps to reduce underflow problems from small probabilities: e.g., vocabulary size of 100,000)
 - If the probabilities look random, the model hasn't learned much
 - Helpful for debugging
- Log Probs = $\text{Log}(\text{Probs})$: Per-token probabilities for analysis
- Log Probs can be used as a level of confidence (i.e., high likely tokens)

Why it matters

- Understanding model confidence: You can see the probability distribution over tokens.
- Useful for advanced applications: You might re-rank or filter tokens yourself based on those probabilities.
- Helps in prompt engineering: By looking at tokens with high or low probabilities, you may identify ambiguous phrasing in your prompt that is leading to inconsistent output

TOP_P (NUCLEUS SAMPLING)

What it is

- An alternative (or complementary) way of controlling randomness. top_p is the **cumulative probability threshold**. The model will consider **only the tokens within the “top” probability mass**.

Why it matters

- When top_p is smaller (e.g., 0.1), the model focuses on **highly probable tokens**, making responses **more conservative and repetitive**.
- When top_p is larger (close to 1.0), the model considers a **wider distribution of potential tokens**, yielding **more diverse outputs**.

PENALTIES (1/2)

Frequency_Penalty

What it is

- A value between -2.0 and 2.0 that penalizes or rewards new tokens based on their frequency so far in the text. A higher positive value will make the model *less likely* to repeat tokens it has already used.

Why it matters

- Reduces redundancy: If the model tends to get repetitive, increasing frequency_penalty discourages repeated tokens.
- Negative values can encourage repetition if that's desired (rare).

Presence_Penalty

What it is

- Similar to frequency_penalty but slightly different in logic. A higher positive value penalizes tokens that have already appeared in the text, *regardless* of how many times they appear.

Why it matters

- Encourages novelty: Even if a word has appeared once, the model is more likely to try something new.
- Helps avoid repeating entire lines or phrases.

PENALTIES: EXAMPLES (2/2)

Frequency_Penalty

Example:

```
json Copy  
  
{  
  "prompt": "Write a paragraph describing an autumn day.",  
  "max_tokens": 60,  
  "frequency_penalty": 1.0  
}
```

- The model will try not to overuse the same words like "leaves" or "trees," leading to more varied vocabulary.

Presence_Penalty

Example:

```
json Copy  
  
{  
  "prompt": "Give me instructions on how to plant a rose garden.",  
  "presence_penalty": 0.5  
}
```

- With a `presence_penalty` of 0.5, the model will try to avoid using the same keywords repeatedly, so it might use synonyms or rephrase steps for variety.

MAX_TOKENS

```
{
  "prompt": "Write a short story about a space explorer who discovers a new planet.",
  "max_tokens": 50
}
```

Example:

Here, we tell the model to generate at most 50 tokens. That might result in only 3–5 sentences of output (depending on the tokens). If you need a **longer story**, you'd increase **max tokens to 200+**

What it is:

- This parameter sets the **maximum number of tokens (word pieces)** that the **model can generate in its response**.

Why it matters:

- Controlling length: If you have a strict limit on how long the output can be (e.g., a tweet-like format), you can set **max tokens** to cap the model's response.
- Efficiency: A **smaller max tokens** will typically be **cheaper and faster**. However, it might **cut off the response mid-sentence** if it's too restrictive.

STOP (OR STOP_SEQUENCES)

What it is

- A string or list of strings the model will stop generating tokens upon encountering. This effectively truncates the output at a certain pattern.

Why it matters

- Truncating the response safely: For example, you might want to stop the model when it starts a new line or sees a certain sentinel token.
- Useful for structured outputs: If you're generating JSON or code, you can define a stop sequence that ends the code block.

Example:

```
json Copy  
  
{  
  "prompt": "List three advantages of renewable energy:\n1.",  
  "stop": ["4."]   
}
```

- The model will stop generating once it begins writing "4." (or hits the token boundary for "4."), ensuring only three advantages are listed.

N AND BEST_OF

What they are

- **n**: The number of completions to generate for each prompt.
- **best_of**: When used, it generates multiple completions server-side but returns only the “best” one according to some criterion (usually likelihood).

Why they matter

- **n** can help you sample multiple possible outputs for a creative or brainstorming scenario in a single request.
- **best_of** is useful to get the highest probability completion from a set, though note it uses more compute (and thus may cost more).

Example:

```
json
{
  "prompt": "Suggest a tagline for an eco-friendly reusable water bottle",
  "n": 3,
  "best_of": 3
}
```

- The API might internally generate 3 completions and return the best one (if you also set `best_of`). If `n=3` without `best_of`, you'll simply get all 3 completions back.

PROMPT ENGINEERING 201 DEMO

- *NOTEBOOK: `PROMPT_ENGINEERING.IPYNB`*



DEMO